

assignment05_SafsafiAchraf

April 17, 2021

Achraf Safsafi

DSC650

Assignment 5

Assignment 5.1 : Classifying movie reviews

The IMDB dataset

Loading the IMDB dataset

```
[1]: import warnings
      warnings.filterwarnings("ignore")
```

```
[2]: from keras.datasets import imdb

      (train_data, train_labels), (test_data, test_labels) = imdb.
      ↳load_data(num_words=10000)
```

```
[3]: print(train_data[0])
```

```
[1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36,
256, 5, 25, 100, 43, 838, 112, 50, 670, 2, 9, 35, 480, 284, 5, 150, 4, 172, 112,
167, 2, 336, 385, 39, 4, 172, 4536, 1111, 17, 546, 38, 13, 447, 4, 192, 50, 16,
6, 147, 2025, 19, 14, 22, 4, 1920, 4613, 469, 4, 22, 71, 87, 12, 16, 43, 530,
38, 76, 15, 13, 1247, 4, 22, 17, 515, 17, 12, 16, 626, 18, 2, 5, 62, 386, 12, 8,
316, 8, 106, 5, 4, 2223, 5244, 16, 480, 66, 3785, 33, 4, 130, 12, 16, 38, 619,
5, 25, 124, 51, 36, 135, 48, 25, 1415, 33, 6, 22, 12, 215, 28, 77, 52, 5, 14,
407, 16, 82, 2, 8, 4, 107, 117, 5952, 15, 256, 4, 2, 7, 3766, 5, 723, 36, 71,
43, 530, 476, 26, 400, 317, 46, 7, 4, 2, 1029, 13, 104, 88, 4, 381, 15, 297, 98,
32, 2071, 56, 26, 141, 6, 194, 7486, 18, 4, 226, 22, 21, 134, 476, 26, 480, 5,
144, 30, 5535, 18, 51, 36, 28, 224, 92, 25, 104, 4, 226, 65, 16, 38, 1334, 88,
12, 16, 283, 5, 16, 4472, 113, 103, 32, 15, 16, 5345, 19, 178, 32]
```

```
[4]: train_labels[0]
```

```
[4]: 1
```

No word index will exceed 10,000

```
[5]: max([max(sequence) for sequence in train_data])
```

```
[5]: 9999
```

Decode one of these reviews back to English words

```
[6]: word_index = imdb.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
decoded_review = ' '.join([reverse_word_index.get(i - 3, '?') for i in
    ↪train_data[0]])
decoded_review
```

```
[6]: "? this film was just brilliant casting location scenery story direction
everyone's really suited the part they played and you could just imagine being
there robert ? is an amazing actor and now the same being director ? father came
from the same scottish island as myself so i loved the fact there was a real
connection with this film the witty remarks throughout the film were great it
was just brilliant so much that i bought the film as soon as it was released for
? and would recommend it to everyone to watch and the fly fishing was amazing
really cried at the end it was so sad and you know what they say if you cry at a
film it must have been good and this definitely was also ? to the two little
boy's that played the ? of norman and paul they were just brilliant children are
often left out of the ? list i think because the stars that play them all grown
up are such a big profile for the whole film but these children are amazing and
should be praised for what they have done don't you think the whole story was so
lovely because it was true and was someone's life after all that was shared with
us all"
```

Preparing the data

Encoding the integer sequences into a binary matrix

```
[7]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)

x_train[0]
```

```
[7]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
[8]: y_train = np.asarray(train_labels).astype('float32')
     y_test = np.asarray(test_labels).astype('float32')
```

Building our network

The model definition

```
[9]: from keras import models
     from keras import layers

     model = models.Sequential()
     model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
     model.add(layers.Dense(16, activation='relu'))
     model.add(layers.Dense(1, activation='sigmoid'))
```

Compiling the model

```
[10]: model.compile(optimizer='rmsprop',
                   loss='binary_crossentropy',
                   metrics=['accuracy'])
```

Configuring the optimizer

```
[11]: from keras import optimizers

     model.compile(optimizer=optimizers.RMSprop(lr=0.001),
                   loss='binary_crossentropy',
                   metrics=['accuracy'])
```

Using custom losses and metrics

```
[12]: from keras import losses
     from keras import metrics

     model.compile(optimizer=optimizers.RMSprop(lr=0.001),
                   loss=losses.binary_crossentropy,
                   metrics=[metrics.binary_accuracy])
```

Validating our approach

Setting aside a validation set

```
[13]: x_val = x_train[:10000]
     partial_x_train = x_train[10000:]

     y_val = y_train[:10000]
     partial_y_train = y_train[10000:]
```

Training our model

```
[14]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 2s 47ms/step - loss: 0.6099 -
binary_accuracy: 0.6801 - val_loss: 0.3972 - val_binary_accuracy: 0.8700

Epoch 2/20

30/30 [=====] - 0s 14ms/step - loss: 0.3336 -
binary_accuracy: 0.9005 - val_loss: 0.3558 - val_binary_accuracy: 0.8578

Epoch 3/20

30/30 [=====] - 0s 12ms/step - loss: 0.2402 -
binary_accuracy: 0.9264 - val_loss: 0.3086 - val_binary_accuracy: 0.8749

Epoch 4/20

30/30 [=====] - 0s 13ms/step - loss: 0.1872 -
binary_accuracy: 0.9401 - val_loss: 0.2730 - val_binary_accuracy: 0.8916

Epoch 5/20

30/30 [=====] - 0s 13ms/step - loss: 0.1453 -
binary_accuracy: 0.9564 - val_loss: 0.3039 - val_binary_accuracy: 0.8803

Epoch 6/20

30/30 [=====] - 0s 12ms/step - loss: 0.1189 -
binary_accuracy: 0.9652 - val_loss: 0.2917 - val_binary_accuracy: 0.8858

Epoch 7/20

30/30 [=====] - 0s 12ms/step - loss: 0.0975 -
binary_accuracy: 0.9723 - val_loss: 0.3244 - val_binary_accuracy: 0.8825

Epoch 8/20

30/30 [=====] - 0s 13ms/step - loss: 0.0818 -
binary_accuracy: 0.9783 - val_loss: 0.3298 - val_binary_accuracy: 0.8825

Epoch 9/20

30/30 [=====] - 0s 13ms/step - loss: 0.0655 -
binary_accuracy: 0.9854 - val_loss: 0.3941 - val_binary_accuracy: 0.8730

Epoch 10/20

30/30 [=====] - 0s 12ms/step - loss: 0.0585 -
binary_accuracy: 0.9856 - val_loss: 0.3770 - val_binary_accuracy: 0.8801

Epoch 11/20

30/30 [=====] - 0s 11ms/step - loss: 0.0458 -
binary_accuracy: 0.9895 - val_loss: 0.4018 - val_binary_accuracy: 0.8756

Epoch 12/20

30/30 [=====] - 0s 12ms/step - loss: 0.0359 -
binary_accuracy: 0.9932 - val_loss: 0.4294 - val_binary_accuracy: 0.8778

Epoch 13/20

30/30 [=====] - 0s 12ms/step - loss: 0.0262 -
binary_accuracy: 0.9956 - val_loss: 0.5375 - val_binary_accuracy: 0.8543

Epoch 14/20

```

30/30 [=====] - 0s 12ms/step - loss: 0.0246 -
binary_accuracy: 0.9962 - val_loss: 0.5165 - val_binary_accuracy: 0.8634
Epoch 15/20
30/30 [=====] - 0s 12ms/step - loss: 0.0180 -
binary_accuracy: 0.9978 - val_loss: 0.5300 - val_binary_accuracy: 0.8672
Epoch 16/20
30/30 [=====] - 0s 12ms/step - loss: 0.0134 -
binary_accuracy: 0.9986 - val_loss: 0.5569 - val_binary_accuracy: 0.8702
Epoch 17/20
30/30 [=====] - 0s 13ms/step - loss: 0.0099 -
binary_accuracy: 0.9995 - val_loss: 0.5906 - val_binary_accuracy: 0.8694
Epoch 18/20
30/30 [=====] - 0s 12ms/step - loss: 0.0077 -
binary_accuracy: 0.9996 - val_loss: 0.6274 - val_binary_accuracy: 0.8658
Epoch 19/20
30/30 [=====] - 0s 11ms/step - loss: 0.0062 -
binary_accuracy: 0.9996 - val_loss: 0.6620 - val_binary_accuracy: 0.8655
Epoch 20/20
30/30 [=====] - 0s 10ms/step - loss: 0.0040 -
binary_accuracy: 0.9999 - val_loss: 0.6933 - val_binary_accuracy: 0.8665

```

```

[15]: history_dict = history.history
      history_dict.keys()

```

```

[15]: dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])

```

Plotting the training and validation loss

```

[16]: import matplotlib.pyplot as plt


acc = history.history['binary_accuracy']
val_acc = history.history['val_binary_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```




Plotting the training and validation accuracy

```
[17]: plt.clf()
acc_values = history_dict['binary_accuracy']
val_acc_values = history_dict['val_binary_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Retraining a model from scratch

```
[18]: model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

Epoch 1/4

49/49 [=====] - 1s 13ms/step - loss: 0.5778 - accuracy: 0.7230

Epoch 2/4

49/49 [=====] - 1s 11ms/step - loss: 0.2876 - accuracy: 0.9055

Epoch 3/4

49/49 [=====] - 0s 9ms/step - loss: 0.2098 - accuracy: 0.9262

Epoch 4/4

```
49/49 [=====] - 0s 8ms/step - loss: 0.1693 - accuracy: 0.9398
782/782 [=====] - 1s 2ms/step - loss: 0.2996 - accuracy: 0.8798
```

```
[19]: results
```

```
[19]: [0.29959574341773987, 0.8797600269317627]
```

Using a trained network to generate predictions on new data

```
[20]: model.predict(x_test)
```

```
[20]: array([[0.16862705],
          [0.9999273 ],
          [0.8069594 ],
          ...,
          [0.07418606],
          [0.06543425],
          [0.4470158 ]], dtype=float32)
```

Assignment 5.2 : Classifying newswires

The Reuters dataset

Loading the Reuters dataset

```
[21]: from keras.datasets import reuters

      (train_data, train_labels), (test_data, test_labels) = reuters.
      ↪load_data(num_words=10000)
```

```
[22]: len(train_data)
```

```
[22]: 8982
```

```
[23]: len(test_data)
```

```
[23]: 2246
```

```
[24]: print(train_data[10])
```

```
[1, 245, 273, 207, 156, 53, 74, 160, 26, 14, 46, 296, 26, 39, 74, 2979, 3554,
14, 46, 4689, 4329, 86, 61, 3499, 4795, 14, 61, 451, 4329, 17, 12]
```

Decoding newswires back to text


```
[25]: word_index = reuters.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

decoded_newswire = ' '.join([reverse_word_index.get(i - 3, '?') for i in
    ↪train_data[0]])
decoded_newswire
```

```
[25]: '??? said as a result of its december acquisition of space co it expects
earnings per share in 1987 of 1 15 to 1 30 dlrs per share up from 70 cts in 1986
the company said pretax net should rise to nine to 10 mln dlrs from six mln dlrs
in 1986 and rental operation revenues to 19 to 22 mln dlrs from 12 5 mln dlrs it
said cash flow per share this year should be 2 50 to three dlrs reuter 3'
```

```
[26]: train_labels[10]
```

```
[26]: 3
```

Preparing the data

Encoding the data

```
[27]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
[28]: def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

one_hot_train_labels = to_one_hot(train_labels)
one_hot_test_labels = to_one_hot(test_labels)
```

```
[29]: from keras.utils.np_utils import to_categorical

one_hot_train_labels = to_categorical(train_labels)
one_hot_test_labels = to_categorical(test_labels)
```

Building our network

Model definition

```
[30]: from keras import models
      from keras import layers

      model = models.Sequential()
      model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
      model.add(layers.Dense(64, activation='relu'))
      model.add(layers.Dense(46, activation='softmax'))
```

Compiling the model

```
[31]: model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

Validating our approach

Setting aside a validation set

```
[32]: x_val = x_train[:1000]
      partial_x_train = x_train[1000:]

      y_val = one_hot_train_labels[:1000]
      partial_y_train = one_hot_train_labels[1000:]
```

Training the model

```
[33]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))
```

Epoch 1/20

16/16 [=====] - 1s 29ms/step - loss: 3.0710 - accuracy: 0.4364 - val_loss: 1.6837 - val_accuracy: 0.6370

Epoch 2/20

16/16 [=====] - 0s 17ms/step - loss: 1.4674 - accuracy: 0.6899 - val_loss: 1.2818 - val_accuracy: 0.7150

Epoch 3/20

16/16 [=====] - 0s 16ms/step - loss: 1.0750 - accuracy: 0.7643 - val_loss: 1.1426 - val_accuracy: 0.7390

Epoch 4/20

16/16 [=====] - 0s 16ms/step - loss: 0.8202 - accuracy: 0.8278 - val_loss: 1.0368 - val_accuracy: 0.7700

Epoch 5/20

16/16 [=====] - 0s 21ms/step - loss: 0.6546 - accuracy: 0.8640 - val_loss: 0.9590 - val_accuracy: 0.7940

Epoch 6/20

```

16/16 [=====] - 0s 17ms/step - loss: 0.5272 - accuracy:
0.8964 - val_loss: 0.9218 - val_accuracy: 0.7990
Epoch 7/20
16/16 [=====] - 0s 15ms/step - loss: 0.4383 - accuracy:
0.9109 - val_loss: 0.9115 - val_accuracy: 0.8080
Epoch 8/20
16/16 [=====] - 0s 16ms/step - loss: 0.3454 - accuracy:
0.9268 - val_loss: 0.9344 - val_accuracy: 0.7910
Epoch 9/20
16/16 [=====] - 0s 15ms/step - loss: 0.2834 - accuracy:
0.9405 - val_loss: 0.9008 - val_accuracy: 0.8040
Epoch 10/20
16/16 [=====] - 0s 15ms/step - loss: 0.2435 - accuracy:
0.9439 - val_loss: 1.0188 - val_accuracy: 0.7780
Epoch 11/20
16/16 [=====] - 0s 15ms/step - loss: 0.2132 - accuracy:
0.9504 - val_loss: 0.9124 - val_accuracy: 0.8040
Epoch 12/20
16/16 [=====] - 0s 14ms/step - loss: 0.1784 - accuracy:
0.9536 - val_loss: 0.9277 - val_accuracy: 0.8010
Epoch 13/20
16/16 [=====] - 0s 15ms/step - loss: 0.1532 - accuracy:
0.9571 - val_loss: 0.9558 - val_accuracy: 0.8130
Epoch 14/20
16/16 [=====] - 0s 15ms/step - loss: 0.1391 - accuracy:
0.9598 - val_loss: 0.9694 - val_accuracy: 0.8060
Epoch 15/20
16/16 [=====] - 0s 17ms/step - loss: 0.1305 - accuracy:
0.9614 - val_loss: 0.9987 - val_accuracy: 0.8050
Epoch 16/20
16/16 [=====] - 0s 17ms/step - loss: 0.1146 - accuracy:
0.9621 - val_loss: 1.0329 - val_accuracy: 0.7930
Epoch 17/20
16/16 [=====] - 0s 17ms/step - loss: 0.1232 - accuracy:
0.9570 - val_loss: 1.0481 - val_accuracy: 0.8010
Epoch 18/20
16/16 [=====] - 0s 15ms/step - loss: 0.1140 - accuracy:
0.9594 - val_loss: 1.1128 - val_accuracy: 0.7920
Epoch 19/20
16/16 [=====] - 0s 15ms/step - loss: 0.1091 - accuracy:
0.9637 - val_loss: 1.1062 - val_accuracy: 0.7900
Epoch 20/20
16/16 [=====] - 0s 17ms/step - loss: 0.1016 - accuracy:
0.9656 - val_loss: 1.0945 - val_accuracy: 0.7950

```

Plotting the training and validation loss


```
[34]: import matplotlib.pyplot as plt

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```




Plotting the training and validation accuracy

```
[35]: plt.clf()
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
```

```
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Retraining a model from scratch

```
[36]: model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(partial_x_train,
        partial_y_train,
        epochs=8,
        batch_size=512,
        validation_data=(x_val, y_val))
results = model.evaluate(x_test, one_hot_test_labels)
```

```

Epoch 1/8
16/16 [=====] - 1s 25ms/step - loss: 3.1525 - accuracy:
0.3970 - val_loss: 1.7930 - val_accuracy: 0.6240
Epoch 2/8
16/16 [=====] - 0s 16ms/step - loss: 1.5477 - accuracy:
0.6829 - val_loss: 1.3402 - val_accuracy: 0.6980
Epoch 3/8
16/16 [=====] - 0s 16ms/step - loss: 1.1257 - accuracy:
0.7537 - val_loss: 1.1730 - val_accuracy: 0.7380
Epoch 4/8
16/16 [=====] - 0s 18ms/step - loss: 0.8850 - accuracy:
0.8123 - val_loss: 1.0582 - val_accuracy: 0.7710
Epoch 5/8
16/16 [=====] - 0s 15ms/step - loss: 0.7280 - accuracy:
0.8466 - val_loss: 0.9824 - val_accuracy: 0.7870
Epoch 6/8
16/16 [=====] - 0s 16ms/step - loss: 0.5632 - accuracy:
0.8823 - val_loss: 0.9400 - val_accuracy: 0.8100
Epoch 7/8
16/16 [=====] - 0s 16ms/step - loss: 0.4505 - accuracy:
0.9103 - val_loss: 0.8879 - val_accuracy: 0.8200
Epoch 8/8
16/16 [=====] - 0s 14ms/step - loss: 0.3563 - accuracy:
0.9276 - val_loss: 0.8959 - val_accuracy: 0.8210
71/71 [=====] - 0s 2ms/step - loss: 0.9783 - accuracy:
0.7872

```

```
[37]: results
```

```
[37]: [0.9783440828323364, 0.7871772050857544]
```

The accuracy reached by the random classifier is about 19%

```

[38]: import copy

test_labels_copy = copy.copy(test_labels)
np.random.shuffle(test_labels_copy)
float(np.sum(np.array(test_labels) == np.array(test_labels_copy))) /
↳ len(test_labels)

```

```
[38]: 0.19545859305431879
```

Generating predictions on new data

Generating predictions for new data

```
[39]: predictions = model.predict(x_test)
```

Each entry in predictions is a vector of length 46:

```
[40]: predictions[0].shape
```

```
[40]: (46,)
```

The coefficients in this vector sum to 1:

```
[41]: np.sum(predictions[0])
```

```
[41]: 0.9999999
```

The largest entry is the predicted class—the class with the highest probability:

```
[42]: np.argmax(predictions[0])
```

```
[42]: 3
```

A different way to handle the labels and the loss

```
[43]: y_train = np.array(train_labels)
      y_test = np.array(test_labels)
```

```
[44]: model.compile(optimizer='rmsprop', loss='sparse_categorical_crossentropy',
      ↪metrics=['acc'])
```

The importance of having sufficiently large intermediate layers

A model with an information bottleneck

```
[45]: model = models.Sequential()
      model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
      model.add(layers.Dense(4, activation='relu'))
      model.add(layers.Dense(46, activation='softmax'))

      model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
      model.fit(partial_x_train,
                partial_y_train,
                epochs=20,
                batch_size=128,
                validation_data=(x_val, y_val))
```

Epoch 1/20

63/63 [=====] - 1s 14ms/step - loss: 3.5529 - accuracy:
0.1590 - val_loss: 2.7013 - val_accuracy: 0.2520

Epoch 2/20

63/63 [=====] - 0s 6ms/step - loss: 2.4127 - accuracy:

0.3419 - val_loss: 1.8286 - val_accuracy: 0.5450
Epoch 3/20
63/63 [=====] - 0s 7ms/step - loss: 1.6928 - accuracy: 0.5567 - val_loss: 1.5804 - val_accuracy: 0.5720
Epoch 4/20
63/63 [=====] - 0s 6ms/step - loss: 1.3713 - accuracy: 0.6072 - val_loss: 1.4630 - val_accuracy: 0.6480
Epoch 5/20
63/63 [=====] - 0s 6ms/step - loss: 1.2216 - accuracy: 0.6859 - val_loss: 1.3901 - val_accuracy: 0.6760
Epoch 6/20
63/63 [=====] - 0s 7ms/step - loss: 1.0505 - accuracy: 0.7205 - val_loss: 1.3382 - val_accuracy: 0.6910
Epoch 7/20
63/63 [=====] - 0s 6ms/step - loss: 0.9709 - accuracy: 0.7533 - val_loss: 1.3063 - val_accuracy: 0.7030
Epoch 8/20
63/63 [=====] - 0s 7ms/step - loss: 0.8349 - accuracy: 0.7954 - val_loss: 1.3540 - val_accuracy: 0.7050
Epoch 9/20
63/63 [=====] - 0s 7ms/step - loss: 0.7643 - accuracy: 0.8046 - val_loss: 1.3584 - val_accuracy: 0.6980
Epoch 10/20
63/63 [=====] - 0s 7ms/step - loss: 0.6983 - accuracy: 0.8252 - val_loss: 1.3388 - val_accuracy: 0.7130
Epoch 11/20
63/63 [=====] - 0s 7ms/step - loss: 0.6577 - accuracy: 0.8301 - val_loss: 1.3590 - val_accuracy: 0.7200
Epoch 12/20
63/63 [=====] - 0s 7ms/step - loss: 0.6226 - accuracy: 0.8393 - val_loss: 1.3927 - val_accuracy: 0.7120
Epoch 13/20
63/63 [=====] - 0s 6ms/step - loss: 0.5623 - accuracy: 0.8580 - val_loss: 1.4391 - val_accuracy: 0.7190
Epoch 14/20
63/63 [=====] - 0s 7ms/step - loss: 0.5325 - accuracy: 0.8610 - val_loss: 1.4704 - val_accuracy: 0.7180
Epoch 15/20
63/63 [=====] - 0s 7ms/step - loss: 0.4871 - accuracy: 0.8758 - val_loss: 1.5119 - val_accuracy: 0.7230
Epoch 16/20
63/63 [=====] - 0s 7ms/step - loss: 0.4781 - accuracy: 0.8766 - val_loss: 1.5572 - val_accuracy: 0.7160
Epoch 17/20
63/63 [=====] - 0s 6ms/step - loss: 0.4555 - accuracy: 0.8790 - val_loss: 1.6022 - val_accuracy: 0.7170
Epoch 18/20
63/63 [=====] - 0s 7ms/step - loss: 0.4328 - accuracy:


```

0.8860 - val_loss: 1.6134 - val_accuracy: 0.7220
Epoch 19/20
63/63 [=====] - 0s 7ms/step - loss: 0.4052 - accuracy:
0.8881 - val_loss: 1.6638 - val_accuracy: 0.7140
Epoch 20/20
63/63 [=====] - 0s 6ms/step - loss: 0.4018 - accuracy:
0.8941 - val_loss: 1.7169 - val_accuracy: 0.7210

```

```
[45]: <tensorflow.python.keras.callbacks.History at 0x7f4a341dc220>
```

Assignment 5.3 : Predicting house prices

The Boston Housing Price dataset

Loading the Boston housing dataset

```
[46]: from keras.datasets import boston_housing

(train_data, train_targets), (test_data, test_targets) = boston_housing.
↳load_data()
```

```
[47]: train_data.shape
```

```
[47]: (404, 13)
```

```
[48]: test_data.shape
```

```
[48]: (102, 13)
```

```
[49]: train_targets
```

```
[49]: array([15.2, 42.3, 50. , 21.1, 17.7, 18.5, 11.3, 15.6, 15.6, 14.4, 12.1,
        17.9, 23.1, 19.9, 15.7,  8.8, 50. , 22.5, 24.1, 27.5, 10.9, 30.8,
        32.9, 24. , 18.5, 13.3, 22.9, 34.7, 16.6, 17.5, 22.3, 16.1, 14.9,
        23.1, 34.9, 25. , 13.9, 13.1, 20.4, 20. , 15.2, 24.7, 22.2, 16.7,
        12.7, 15.6, 18.4, 21. , 30.1, 15.1, 18.7,  9.6, 31.5, 24.8, 19.1,
        22. , 14.5, 11. , 32. , 29.4, 20.3, 24.4, 14.6, 19.5, 14.1, 14.3,
        15.6, 10.5,  6.3, 19.3, 19.3, 13.4, 36.4, 17.8, 13.5, 16.5,  8.3,
        14.3, 16. , 13.4, 28.6, 43.5, 20.2, 22. , 23. , 20.7, 12.5, 48.5,
        14.6, 13.4, 23.7, 50. , 21.7, 39.8, 38.7, 22.2, 34.9, 22.5, 31.1,
        28.7, 46. , 41.7, 21. , 26.6, 15. , 24.4, 13.3, 21.2, 11.7, 21.7,
        19.4, 50. , 22.8, 19.7, 24.7, 36.2, 14.2, 18.9, 18.3, 20.6, 24.6,
        18.2,  8.7, 44. , 10.4, 13.2, 21.2, 37. , 30.7, 22.9, 20. , 19.3,
        31.7, 32. , 23.1, 18.8, 10.9, 50. , 19.6,  5. , 14.4, 19.8, 13.8,
        19.6, 23.9, 24.5, 25. , 19.9, 17.2, 24.6, 13.5, 26.6, 21.4, 11.9,
        22.6, 19.6,  8.5, 23.7, 23.1, 22.4, 20.5, 23.6, 18.4, 35.2, 23.1,
```

```

27.9, 20.6, 23.7, 28. , 13.6, 27.1, 23.6, 20.6, 18.2, 21.7, 17.1,
 8.4, 25.3, 13.8, 22.2, 18.4, 20.7, 31.6, 30.5, 20.3,  8.8, 19.2,
19.4, 23.1, 23. , 14.8, 48.8, 22.6, 33.4, 21.1, 13.6, 32.2, 13.1,
23.4, 18.9, 23.9, 11.8, 23.3, 22.8, 19.6, 16.7, 13.4, 22.2, 20.4,
21.8, 26.4, 14.9, 24.1, 23.8, 12.3, 29.1, 21. , 19.5, 23.3, 23.8,
17.8, 11.5, 21.7, 19.9, 25. , 33.4, 28.5, 21.4, 24.3, 27.5, 33.1,
16.2, 23.3, 48.3, 22.9, 22.8, 13.1, 12.7, 22.6, 15. , 15.3, 10.5,
24. , 18.5, 21.7, 19.5, 33.2, 23.2,  5. , 19.1, 12.7, 22.3, 10.2,
13.9, 16.3, 17. , 20.1, 29.9, 17.2, 37.3, 45.4, 17.8, 23.2, 29. ,
22. , 18. , 17.4, 34.6, 20.1, 25. , 15.6, 24.8, 28.2, 21.2, 21.4,
23.8, 31. , 26.2, 17.4, 37.9, 17.5, 20. ,  8.3, 23.9,  8.4, 13.8,
 7.2, 11.7, 17.1, 21.6, 50. , 16.1, 20.4, 20.6, 21.4, 20.6, 36.5,
 8.5, 24.8, 10.8, 21.9, 17.3, 18.9, 36.2, 14.9, 18.2, 33.3, 21.8,
19.7, 31.6, 24.8, 19.4, 22.8,  7.5, 44.8, 16.8, 18.7, 50. , 50. ,
19.5, 20.1, 50. , 17.2, 20.8, 19.3, 41.3, 20.4, 20.5, 13.8, 16.5,
23.9, 20.6, 31.5, 23.3, 16.8, 14. , 33.8, 36.1, 12.8, 18.3, 18.7,
19.1, 29. , 30.1, 50. , 50. , 22. , 11.9, 37.6, 50. , 22.7, 20.8,
23.5, 27.9, 50. , 19.3, 23.9, 22.6, 15.2, 21.7, 19.2, 43.8, 20.3,
33.2, 19.9, 22.5, 32.7, 22. , 17.1, 19. , 15. , 16.1, 25.1, 23.7,
28.7, 37.2, 22.6, 16.4, 25. , 29.8, 22.1, 17.4, 18.1, 30.3, 17.5,
24.7, 12.6, 26.5, 28.7, 13.3, 10.4, 24.4, 23. , 20. , 17.8,  7. ,
11.8, 24.4, 13.8, 19.4, 25.2, 19.4, 19.4, 29.1])

```

Preparing the data

Normalizing the data

```

[50]: mean = train_data.mean(axis=0)
      train_data -= mean
      std = train_data.std(axis=0)
      train_data /= std

      test_data -= mean
      test_data /= std

```

Building our network

Model definition

```

[51]: from keras import models
      from keras import layers

      def build_model():
          model = models.Sequential()
          model.add(layers.Dense(64, activation='relu',
                                input_shape=(train_data.shape[1],)))
          model.add(layers.Dense(64, activation='relu'))
          model.add(layers.Dense(1))
          model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])

```

```
return model
```

Validating our approach using K-fold validation

K-fold validation

```
[52]: import numpy as np

k = 4
num_val_samples = len(train_data) // k
num_epochs = 100
all_scores = []
for i in range(k):
    print('processing fold #', i)

    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]

    partial_train_data = np.concatenate(
        [train_data[:i * num_val_samples],
         train_data[(i + 1) * num_val_samples:]],
        axis=0)
    partial_train_targets = np.concatenate(
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]],
        axis=0)

    model = build_model()
    model.fit(partial_train_data, partial_train_targets,
              epochs=num_epochs, batch_size=1, verbose=0)

    val_mse, val_mae = model.evaluate(val_data, val_targets, verbose=0)
    all_scores.append(val_mae)
```

```
processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3
```

```
[53]: all_scores
```

```
[53]: [2.2719955444335938, 2.6824657917022705, 2.980651617050171, 2.458961248397827]
```

```
[54]: np.mean(all_scores)
```

```
[54]: 2.5985185503959656
```

Saving the validation logs at each fold

```
[55]: num_epochs = 500
      all_mae_histories = []
      for i in range(k):
          print('processing fold #', i)
          val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
          val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]

          partial_train_data = np.concatenate(
              [train_data[:i * num_val_samples],
               train_data[(i + 1) * num_val_samples:]],
              axis=0)
          partial_train_targets = np.concatenate(
              [train_targets[:i * num_val_samples],
               train_targets[(i + 1) * num_val_samples:]],
              axis=0)

          model = build_model()
          history = model.fit(partial_train_data, partial_train_targets,
                              validation_data=(val_data, val_targets),
                              epochs=num_epochs, batch_size=1, verbose=0)
          mae_history = history.history['val_mae']
          all_mae_histories.append(mae_history)
```

```
processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3
```


Building the history of successive mean K-fold validation scores

```
[56]: average_mae_history = [
      np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]
```

Plotting validation scores

```
[57]: import matplotlib.pyplot as plt

      plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
      plt.xlabel('Epochs')
      plt.ylabel('Validation MAE')
      plt.show()
```




Plotting validation scores, excluding the first 10 data points

```
[58]: def smooth_curve(points, factor=0.9):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

smooth_mae_history = smooth_curve(average_mae_history[10:])

plt.plot(range(1, len(smooth_mae_history) + 1), smooth_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()
```



Training the final model

```
[59]: model = build_model()
      model.fit(train_data, train_targets,
                epochs=80, batch_size=16, verbose=0)
      test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)
```

```
4/4 [=====] - 0s 1ms/step - loss: 18.5483 - mae: 2.6491
```

```
[60]: test_mae_score
```

```
[60]: 2.6490962505340576
```